

Wildfire Detection with Deep Convolutional Neural Network by Transfer Learning

Sophia Hart, M.Sc., M.Eng.

sophia.t.hart@gmail.com

A project for my Data Science Certificate at UC Berkeley Extension

Agenda

Problem and Goal

Computer Vision using Transfer Learning

Review Jupyter Notebook

Results

Conclusion

Future Work

Q&A

Problem

Increase of CO₂ accumulation → Climate Change:

Temperature rise

Wildfires

Drought

Less snow

Intense hurricanes

Arctic ice shrinkage

Sea level rise

Animals endangered ...

Goal

Wildfires Detection Technologies:

Terrestrial camera systems

Satellites

Drones


IoT

Many involve images


Goal: Train model for computer vision using deep learning

Fundamental of Convolutional Neural Networks (CNN)


Neural network
with
fully connected
layers



CNN utilizes receptive field
Neuron not connected to every
neuron from previous layer



Convolution



Kernel for CNN

Sliding window

1	2	3	1	3	5
2	2	5	4	2	5
0	6	9	6	2	2
2	0	1	9	4	0
5	5	4	6	7	6
6	1	3	7	1	5

convolution

*

1	0	-1
1	0	-1
1	0	-1

=

-14			

Feature map

Filter or Kernel


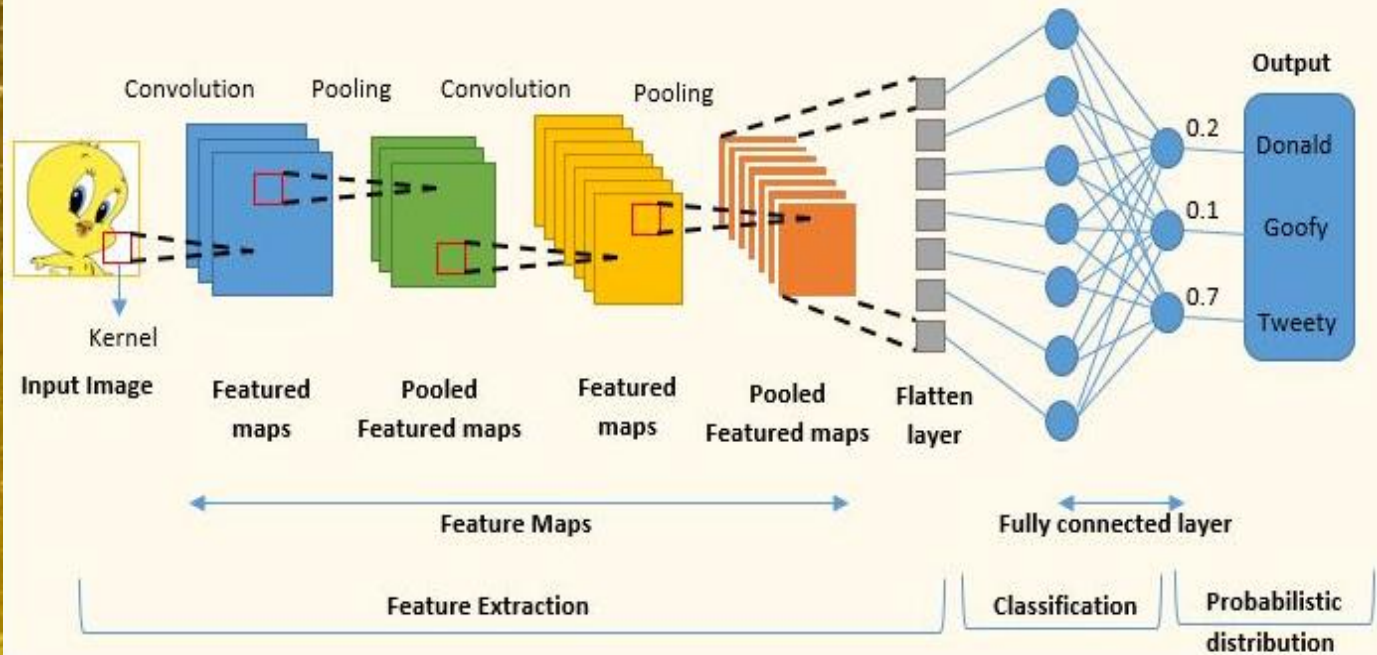
1×1	2×0	3×-1
2×1	2×0	5×-1
0×1	6×0	9×-1

Add up.

$$(1 + 0 - 3) + (2 + 0 - 5) + (0 + 0 - 9) = -14$$

CNN Architecture

A Typical Convolutional Neural Network (CNN)



Goal: Learn the weights/parameters for these kernels in the convolution layers.

Computer Vision using Transfer Learning

Not necessary to reinvent the wheel.

Improved learning of a new task from a related task that's been learned.

If there's not enough training data, good idea to reuse lower layers of pretrained model.

ImageNet Large Scale Visual Recognition Challenge (ILSVRC)

>15 million labeled images, 1000 category of objects

Took 62,000 GPU hours to train (?)

Question: Is model trained for classifying objects applicable for classifying wildfire?

Computer Vision using Transfer Learning

Keras has access to Pretrained models

Pretrained models types:

- AlexNet
- VGG
- Inception (Google)
- ResNet (Microsoft)
- EfficientNet (Google)

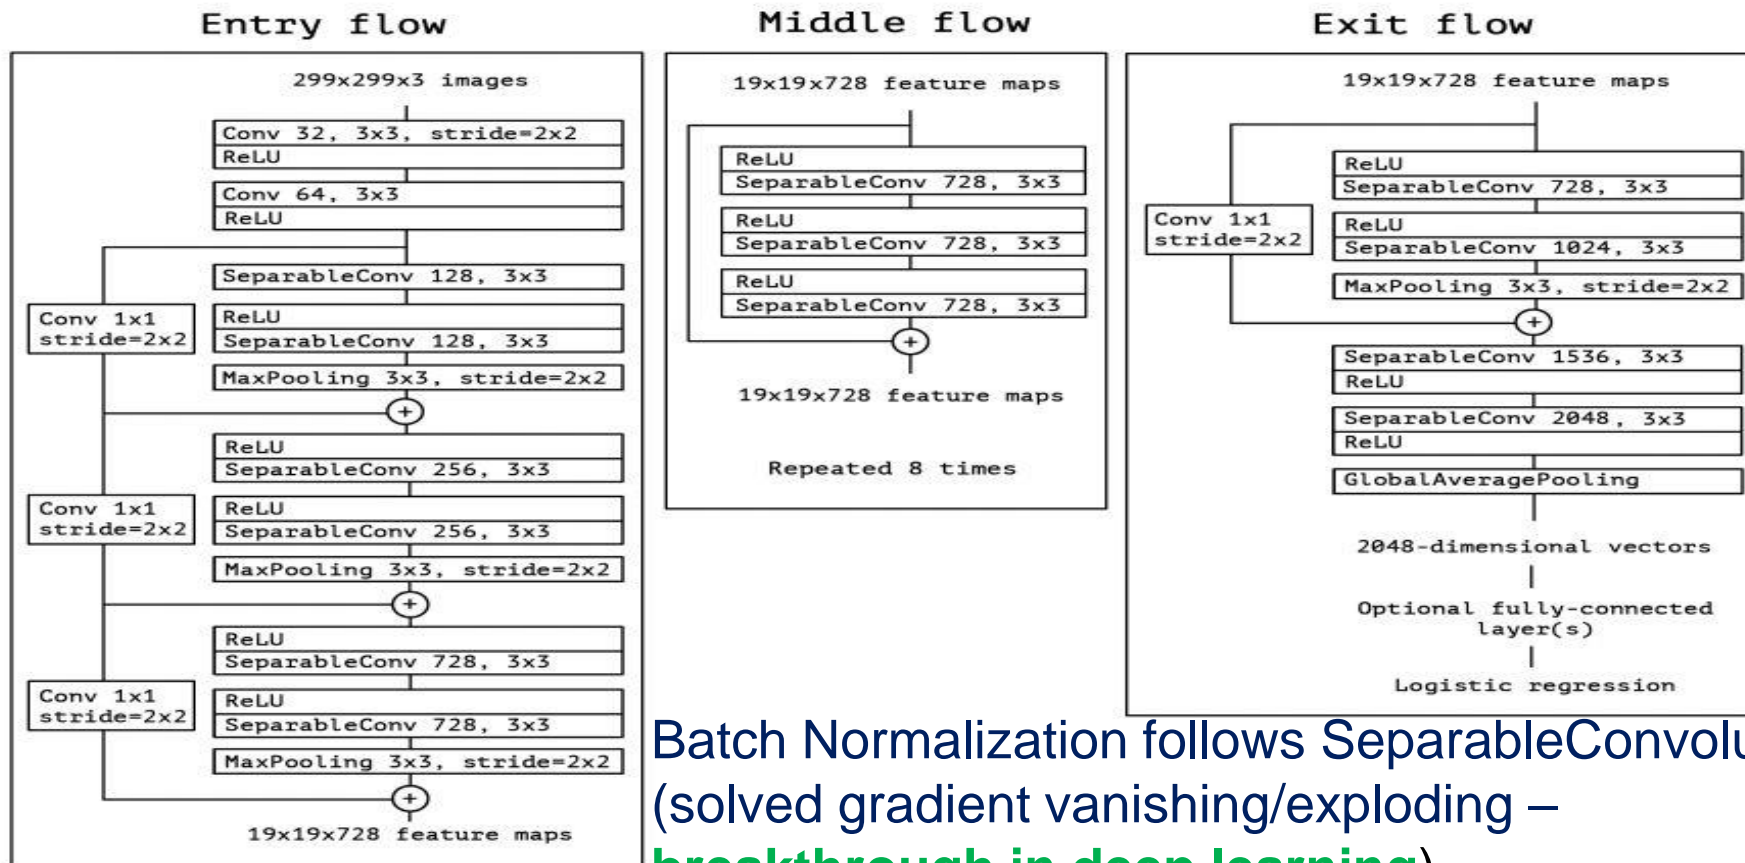
Comparison					
Network	Year	Salient Feature	top5 accuracy	Parameters	FLOP
AlexNet	2012	Deeper	84.70%	62M	1.5B
VGGNet	2014	Fixed-size kernels	92.30%	138M	19.6B
Inception	2014	Wider - Parallel kernels	93.30%	6.4M	2B
ResNet-152	2015	Shortcut connections	95.51%	60.3M	11B

My base model - Xception (Extreme Inception)

Proposed in 2016 by François Chollet, author of Keras

Less memory, exceptional accuracy

Xception Architecture



François Chollet

Batch Normalization follows SeparableConvolution
(solved gradient vanishing/exploding –
breakthrough in deep learning)

Dataset

Forest images

50 no fire

50 with fire

Data Augmentation (flips) 400 images total

Keep 3 channels (color important feature of fire)

Examples of Fire and No-fire Images

Fire



No fire



Challenge of Computer Vision for Wildfire

Fire



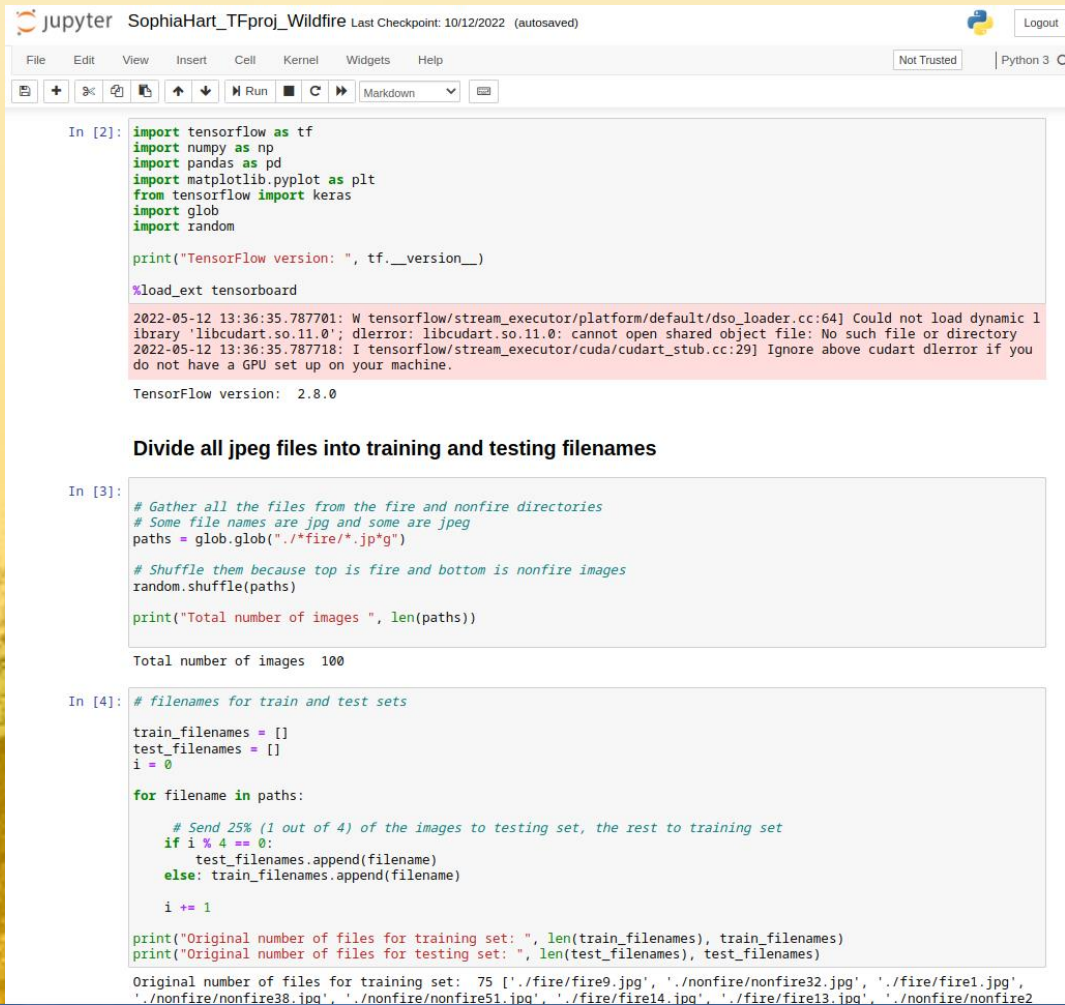
No-fire



Both flame and sunset
or Autumn leaves
appear orange-red

Smoke and cloud
appear similar

Review Jupyter Notebook



Jupyter SophiaHart_TFproj_Wildfire Last Checkpoint: 10/12/2022 (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3

In [2]:

```
import tensorflow as tf
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from tensorflow import keras
import glob
import random

print("TensorFlow version: ", tf.__version__)

%load_ext tensorboard
```

2022-05-12 13:36:35.787701: W tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load dynamic library 'libcudart.so.11.0'; dlderror: libcudart.so.11.0: cannot open shared object file: No such file or directory
2022-05-12 13:36:35.787718: I tensorflow/stream_executor/cuda/cudart_stub.cc:29] Ignore above cudart dlerror if you do not have a GPU set up on your machine.

TensorFlow version: 2.8.0

Divide all jpeg files into training and testing filenames

In [3]:

```
# Gather all the files from the fire and nonfire directories
# Some file names are jpg and some are jpeg
paths = glob.glob("./fire/*.jpg")

# Shuffle them because top is fire and bottom is nonfire images
random.shuffle(paths)

print("Total number of images ", len(paths))
```

Total number of images 100

In [4]:

```
# filenames for train and test sets

train_filenames = []
test_filenames = []
i = 0

for filename in paths:

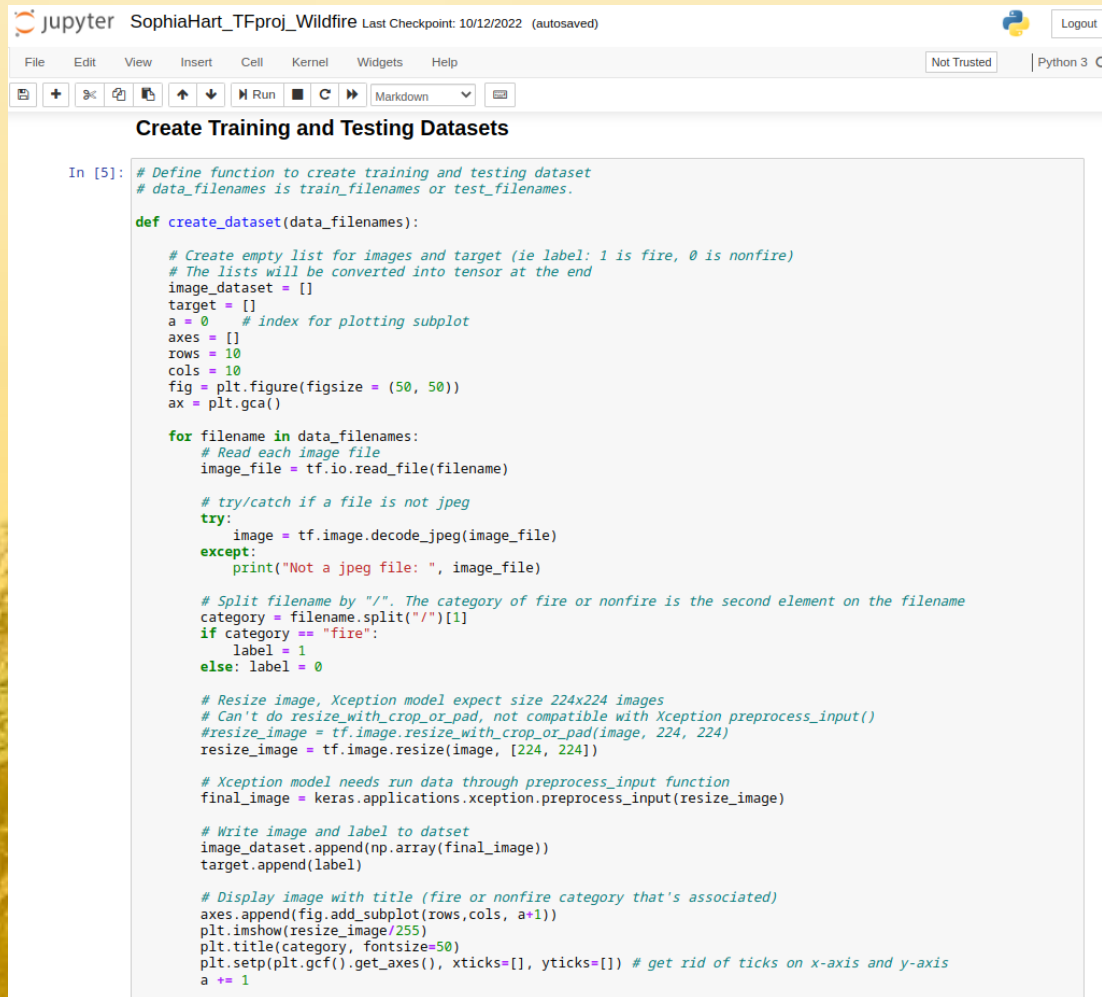
    # Send 25% (1 out of 4) of the images to testing set, the rest to training set
    if i % 4 == 0:
        test_filenames.append(filename)
    else: train_filenames.append(filename)

    i += 1

print("Original number of files for training set: ", len(train_filenames), train_filenames)
print("Original number of files for testing set: ", len(test_filenames), test_filenames)
```

Original number of files for training set: 75 ['./fire/fire9.jpg', './nonfire/nonfire32.jpg', './fire/fire1.jpg', './nonfire/nonfire38.jpg', './nonfire/nonfire51.jpg', './fire/fire14.jpg', './fire/fire13.jpg', './nonfire/nonfire2

Review Jupyter Notebook



The screenshot shows a Jupyter Notebook interface. At the top, the header bar includes the Jupyter logo, the username 'SophiaHart_TFproj_Wildfire', the last checkpoint information 'Last Checkpoint: 10/12/2022 (autosaved)', and a 'Logout' button. Below the header is a menu bar with 'File', 'Edit', 'View', 'Insert', 'Cell', 'Kernel', 'Widgets', and 'Help'. To the right of the menu bar are 'Not Trusted' and 'Python 3 C' indicators. The main toolbar contains icons for file operations, running, and markdown. The notebook content area has a title 'Create Training and Testing Datasets' and a code cell labeled 'In [5]:'. The code cell contains Python code for creating a dataset, including comments and function definitions.

```
In [5]: # Define function to create training and testing dataset
# data_filenames is train_filenames or test_filenames.

def create_dataset(data_filenames):

    # Create empty list for images and target (ie label: 1 is fire, 0 is nonfire)
    # The lists will be converted into tensor at the end
    image_dataset = []
    target = []
    a = 0 # index for plotting subplot
    axes = []
    rows = 10
    cols = 10
    fig = plt.figure(figsize = (50, 50))
    ax = plt.gca()

    for filename in data_filenames:
        # Read each image file
        image_file = tf.io.read_file(filename)

        # try/catch if a file is not jpeg
        try:
            image = tf.image.decode_jpeg(image_file)
        except:
            print("Not a jpeg file: ", image_file)

        # Split filename by "/". The category of fire or nonfire is the second element on the filename
        category = filename.split("/")[1]
        if category == "fire":
            label = 1
        else: label = 0

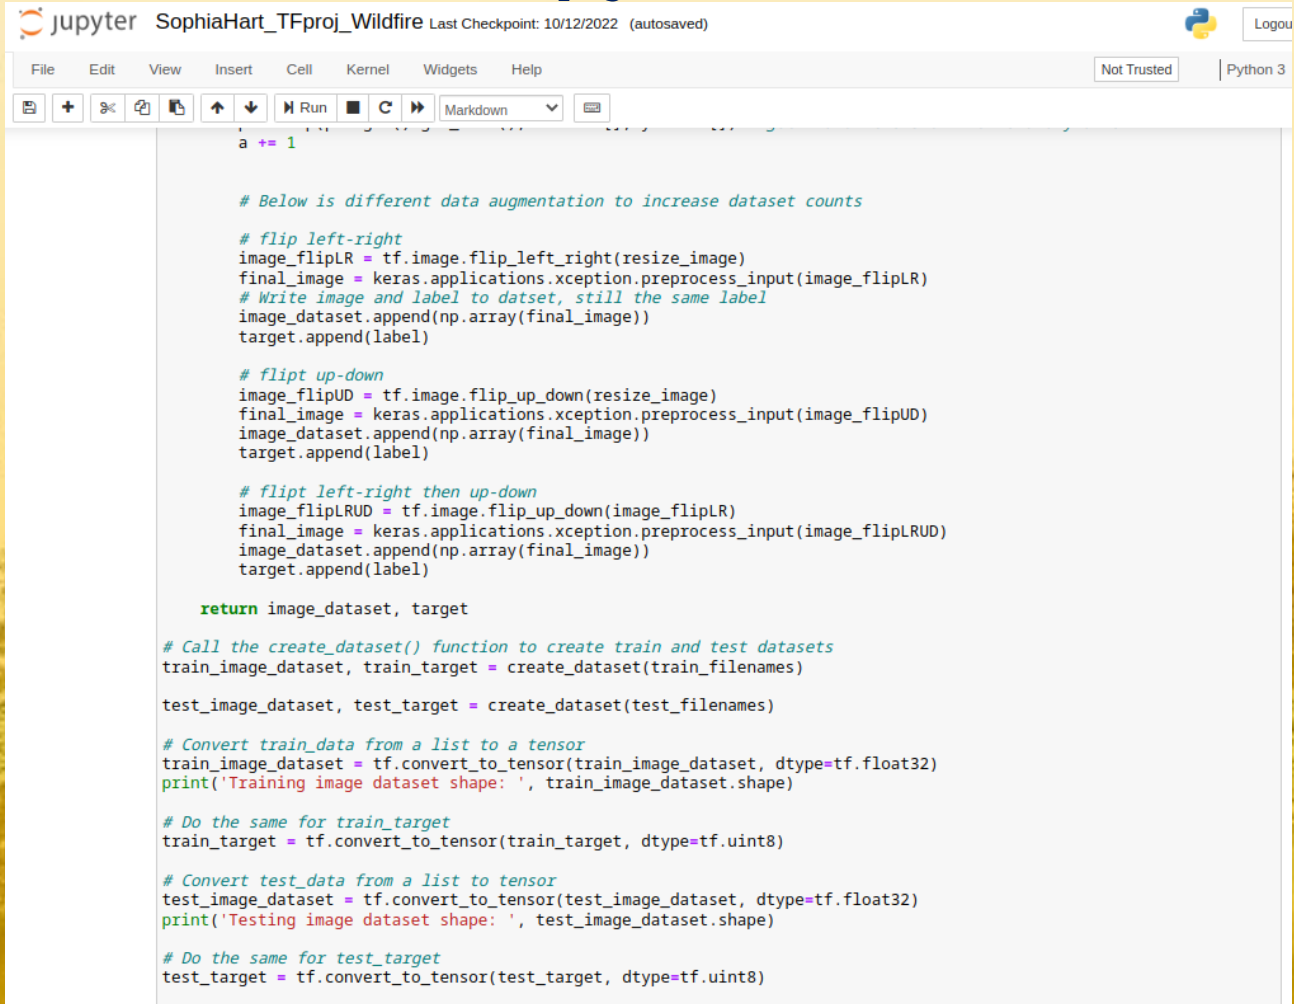
        # Resize image, Xception model expect size 224x224 images
        # Can't do resize_with_crop_or_pad, not compatible with Xception preprocess_input()
        #resize_image = tf.image.resize_with_crop_or_pad(image, 224, 224)
        resize_image = tf.image.resize(image, [224, 224])

        # Xception model needs run data through preprocess_input function
        final_image = keras.applications.xception.preprocess_input(resize_image)

        # Write image and label to dataset
        image_dataset.append(np.array(final_image))
        target.append(label)

        # Display image with title (fire or nonfire category that's associated)
        axes.append(fig.add_subplot(rows,cols, a+1))
        plt.imshow(resize_image/255)
        plt.title(category, fontsize=50)
        plt.setp(plt.gcf().get_axes(), xticks=[], yticks=[]) # get rid of ticks on x-axis and y-axis
        a += 1
```

Review Jupyter Notebook



The screenshot displays a Jupyter Notebook window. The top bar shows the Jupyter logo, the filename 'SophiaHart_TFproj_Wildfire', the last checkpoint '10/12/2022 (autosaved)', a Python 3 logo, and a 'Logout' button. Below the top bar is a menu bar with 'File', 'Edit', 'View', 'Insert', 'Cell', 'Kernel', 'Widgets', and 'Help'. A toolbar contains icons for file operations, a 'Run' button, and a dropdown menu set to 'Markdown'. The main area shows a code cell with the following Python code:

```
a += 1

# Below is different data augmentation to increase dataset counts

# flip left-right
image_flipLR = tf.image.flip_left_right(resize_image)
final_image = keras.applications.xception.preprocess_input(image_flipLR)
# Write image and label to dataset, still the same label
image_dataset.append(np.array(final_image))
target.append(label)

# flitp up-down
image_flipUD = tf.image.flip_up_down(resize_image)
final_image = keras.applications.xception.preprocess_input(image_flipUD)
image_dataset.append(np.array(final_image))
target.append(label)

# flitp left-right then up-down
image_flipLRUD = tf.image.flip_up_down(image_flipLR)
final_image = keras.applications.xception.preprocess_input(image_flipLRUD)
image_dataset.append(np.array(final_image))
target.append(label)

return image_dataset, target

# Call the create_dataset() function to create train and test datasets
train_image_dataset, train_target = create_dataset(train_filenames)

test_image_dataset, test_target = create_dataset(test_filenames)

# Convert train_data from a list to a tensor
train_image_dataset = tf.convert_to_tensor(train_image_dataset, dtype=tf.float32)
print('Training image dataset shape: ', train_image_dataset.shape)

# Do the same for train_target
train_target = tf.convert_to_tensor(train_target, dtype=tf.uint8)

# Convert test_data from a list to tensor
test_image_dataset = tf.convert_to_tensor(test_image_dataset, dtype=tf.float32)
print('Testing image dataset shape: ', test_image_dataset.shape)

# Do the same for test_target
test_target = tf.convert_to_tensor(test_target, dtype=tf.uint8)
```


Training Dataset (part 1 of 3)

Training image dataset shape: (300, 224, 224, 3)

Testing image dataset shape: (100, 224, 224, 3)



Training Dataset (part 2 of 3)



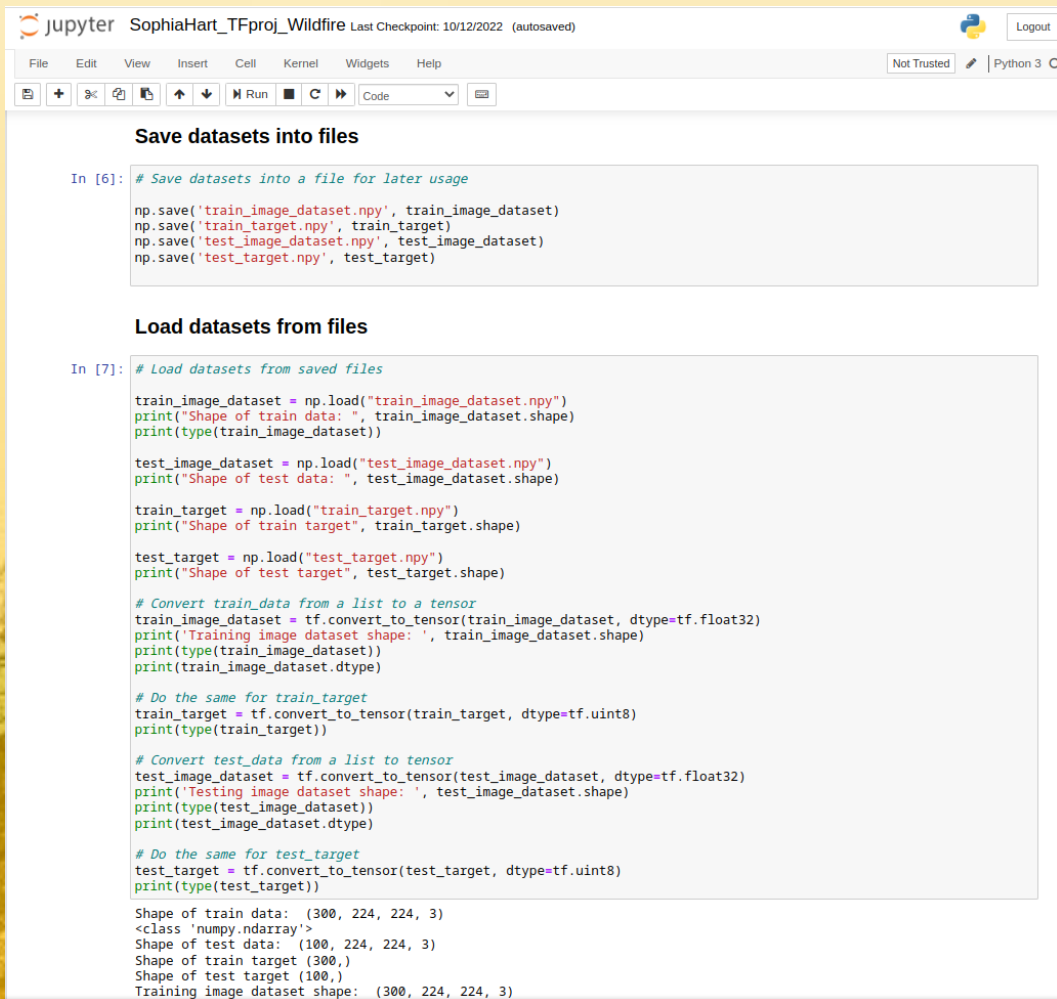
Training Dataset (part 3 of 3)



Test Dataset



Review Jupyter Notebook



The screenshot displays a Jupyter Notebook interface. At the top, the header shows the Jupyter logo, the notebook name 'SophiaHart_TFproj_Wildfire', the last checkpoint '10/12/2022 (autosaved)', and a 'Logout' button. Below the header is a menu bar with 'File', 'Edit', 'View', 'Insert', 'Cell', 'Kernel', 'Widgets', and 'Help'. To the right of the menu bar is a 'Not Trusted' warning and 'Python 3' version information. The main area contains two code cells. The first cell, titled 'Save datasets into files', contains code to save four numpy arrays to files. The second cell, titled 'Load datasets from files', contains code to load these arrays back, convert them to tensors, and print their shapes and dtypes. The output of the second cell is visible at the bottom of the notebook.

```
In [6]: # Save datasets into a file for later usage

np.save('train_image_dataset.npy', train_image_dataset)
np.save('train_target.npy', train_target)
np.save('test_image_dataset.npy', test_image_dataset)
np.save('test_target.npy', test_target)
```

```
In [7]: # Load datasets from saved files

train_image_dataset = np.load("train_image_dataset.npy")
print("Shape of train data: ", train_image_dataset.shape)
print(type(train_image_dataset))

test_image_dataset = np.load("test_image_dataset.npy")
print("Shape of test data: ", test_image_dataset.shape)

train_target = np.load("train_target.npy")
print("Shape of train target", train_target.shape)

test_target = np.load("test_target.npy")
print("Shape of test target", test_target.shape)

# Convert train_data from a list to a tensor
train_image_dataset = tf.convert_to_tensor(train_image_dataset, dtype=tf.float32)
print('Training image dataset shape: ', train_image_dataset.shape)
print(type(train_image_dataset))
print(train_image_dataset.dtype)

# Do the same for train_target
train_target = tf.convert_to_tensor(train_target, dtype=tf.uint8)
print(type(train_target))

# Convert test_data from a list to tensor
test_image_dataset = tf.convert_to_tensor(test_image_dataset, dtype=tf.float32)
print('Testing image dataset shape: ', test_image_dataset.shape)
print(type(test_image_dataset))
print(test_image_dataset.dtype)

# Do the same for test_target
test_target = tf.convert_to_tensor(test_target, dtype=tf.uint8)
print(type(test_target))

Shape of train data: (300, 224, 224, 3)
<class 'numpy.ndarray'>
Shape of test data: (100, 224, 224, 3)
Shape of train target (300,)
Shape of test target (100,)
Training image dataset shape: (300, 224, 224, 3)
```

Review Jupyter Notebook

Jupyter SophiaHart_TFproj_Wildfire Last Checkpoint: 10/12/2022 (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3

Pretrained Model for Transfer Learning

Using Xception model from ImageNet

In [8]: *# Use Xception as the base model*

```
base_model = keras.applications.xception.Xception(weights = 'imagenet', include_top=False)
base_model.summary()
```

Layer (type)	Output Shape	Param #	Connected to
base_model (Xception)	(None, None, None, 1280)	20,861,480	
block14_sepconv1_act (Activation)	(None, None, None, 1536)	0	['block14_sepconv1_bn[0][0]']
block14_sepconv2 (SeparableConv2D)	(None, None, None, 2048)	3159552	['block14_sepconv1_act[0][0]']
block14_sepconv2_bn (BatchNormalization)	(None, None, None, 2048)	8192	['block14_sepconv2[0][0]']
block14_sepconv2_act (Activation)	(None, None, None, 2048)	0	['block14_sepconv2_bn[0][0]']

=====

Total params: 20,861,480
Trainable params: 20,806,952
Non-trainable params: 54,528

In [9]: *# Define model*

```
n_classes = 2 # fire and nonfire
# Output from the base_model
avg = keras.layers.GlobalAveragePooling2D()(base_model.output)
# Only two classes for softmax
output = keras.layers.Dense(n_classes, activation="softmax")(avg)
# Combine base_model and softmax for our model
model = keras.Model(inputs = base_model.input, outputs = output)
# Trainable variables
print("Trainable variables ", len(model.trainable_variables))
model.summary()
```

Trainable variables 156
Model: "model"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, None, None, 0)]	0	[]

Review Jupyter Notebook

```
jupyter SophiaHart_TFproj_Wildfire Last Checkpoint: 10/12/2022 (autosaved) Logout
File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3 C
os.makedirs(logdir, exist_ok=True)

In [11]: # Remove previous event files
!rm -rf ./logs/train/event*

In [12]: tensorboard_callback = keras.callbacks.TensorBoard(logdir)

In [13]: # Freeze the weights of the base model, ie not trainable

for layer in base_model.layers:
    layer.trainable = False

# Compile model.
# Have to use sparse_categorical_crossentropy for compile to work
optimizer = keras.optimizers.SGD(lr=0.01, momentum=0.9, decay=0.01)

model.compile(loss = "sparse_categorical_crossentropy",
              optimizer = optimizer,
              metrics = ['accuracy'])

# Train model
history = model.fit(
    train_image_dataset,
    train_target,
    epochs = 5,
    callbacks = [tensorboard_callback])

/home/shart/schoolwork/TensorFlow_Virtual/tensorflow-dev/lib/python3.7/site-packages/keras/optimizer_v2/gradient_descent.py:102: UserWarning: The 'lr' argument is deprecated, use 'learning_rate' instead.
  super(SGD, self).__init__(name, **kwargs)

Epoch 1/5
10/10 [=====] - 25s 2s/step - loss: 0.5278 - accuracy: 0.7300
Epoch 2/5
10/10 [=====] - 23s 2s/step - loss: 0.1841 - accuracy: 0.9233
Epoch 3/5
10/10 [=====] - 21s 2s/step - loss: 0.1100 - accuracy: 0.9600
Epoch 4/5
10/10 [=====] - 22s 2s/step - loss: 0.0827 - accuracy: 0.9733
Epoch 5/5
10/10 [=====] - 22s 2s/step - loss: 0.0661 - accuracy: 0.9900

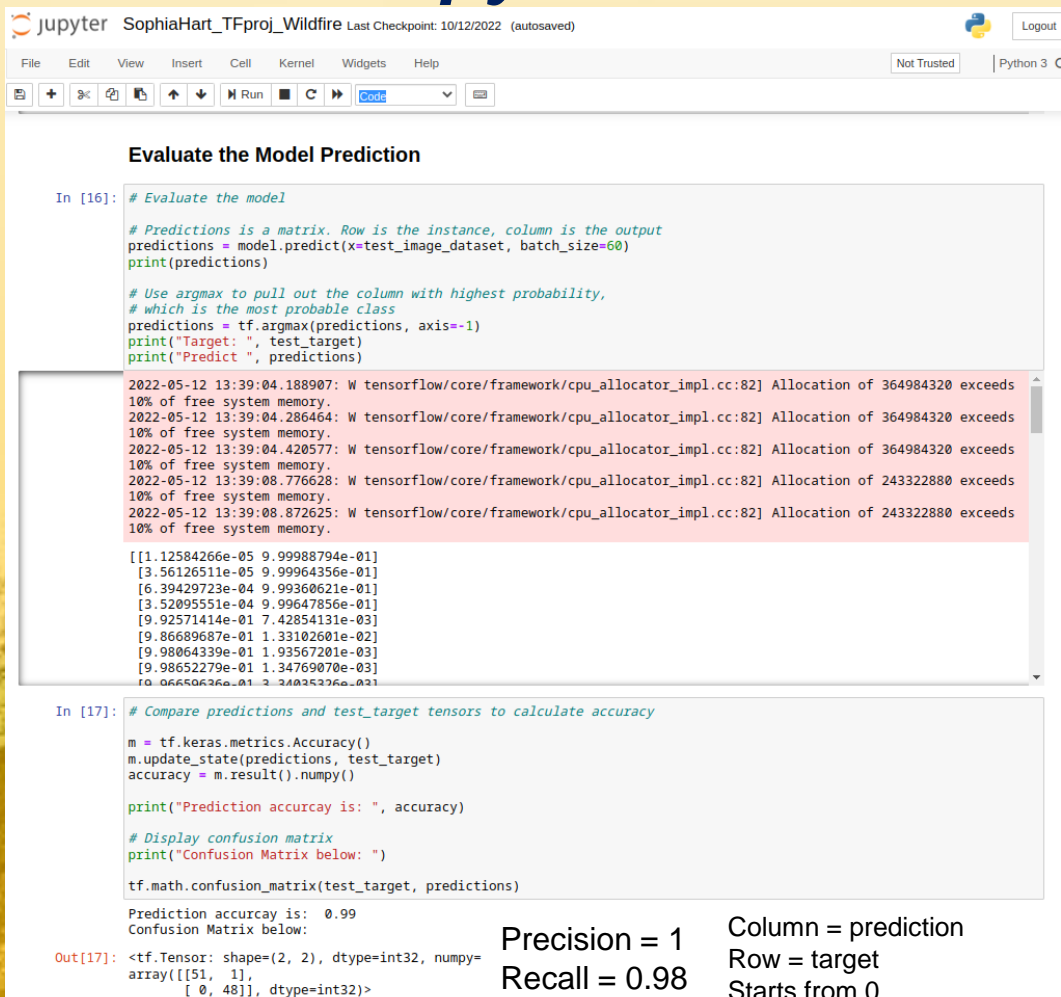
TensorBoard

In [14]: !ls -l ./logs/train

total 404
-rw-rw-r-- 1 shart shart 413510 May 12 13:39 events.out.tfevents.1652387830.ix.11548.0.v2

In [15]: # Show accuracy and loss with each epoch in TensorBoard
%tensorboard --logdir ./logs/train
```

Review Jupyter Notebook



The screenshot shows a Jupyter Notebook interface with the title 'SophiaHart_TFproj_Wildfire'. The interface includes a top bar with the Jupyter logo, the notebook title, and a 'Logout' button. Below the top bar is a menu bar with options: File, Edit, View, Insert, Cell, Kernel, Widgets, and Help. A toolbar with various icons for cell manipulation and execution is also present. The notebook content is divided into two cells. The first cell, titled 'Evaluate the Model Prediction', contains code to evaluate the model's predictions. The second cell, titled 'Compare predictions and test_target tensors to calculate accuracy', contains code to calculate the accuracy and display the confusion matrix. The output of the first cell shows a list of predictions and a list of target labels. The output of the second cell shows the prediction accuracy and the confusion matrix.

```
jupyter SophiaHart_TFproj_Wildfire Last Checkpoint: 10/12/2022 (autosaved) Logout
```

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3 C

Evaluate the Model Prediction

```
In [16]: # Evaluate the model

# Predictions is a matrix. Row is the instance, column is the output
predictions = model.predict(x=test_image_dataset, batch_size=60)
print(predictions)

# Use argmax to pull out the column with highest probability,
# which is the most probable class
predictions = tf.argmax(predictions, axis=-1)
print("Target: ", test_target)
print("Predict ", predictions)
```

```
2022-05-12 13:39:04.188907: W tensorflow/core/framework/cpu_allocator_impl.cc:82] Allocation of 364984320 exceeds
10% of free system memory.
2022-05-12 13:39:04.286464: W tensorflow/core/framework/cpu_allocator_impl.cc:82] Allocation of 364984320 exceeds
10% of free system memory.
2022-05-12 13:39:04.420577: W tensorflow/core/framework/cpu_allocator_impl.cc:82] Allocation of 364984320 exceeds
10% of free system memory.
2022-05-12 13:39:08.776628: W tensorflow/core/framework/cpu_allocator_impl.cc:82] Allocation of 243322880 exceeds
10% of free system memory.
2022-05-12 13:39:08.872625: W tensorflow/core/framework/cpu_allocator_impl.cc:82] Allocation of 243322880 exceeds
10% of free system memory.

[[1.12584266e-05 9.99988794e-01]
 [3.56126511e-05 9.99964356e-01]
 [6.39429723e-04 9.99360621e-01]
 [3.52095551e-04 9.99647856e-01]
 [9.92571414e-01 7.42854131e-03]
 [9.86689687e-01 1.33102601e-02]
 [9.98064339e-01 1.93567201e-03]
 [9.98652279e-01 1.34769070e-03]
 [0.06650636e-01 3.24025276e-02]]
```

```
In [17]: # Compare predictions and test_target tensors to calculate accuracy

m = tf.keras.metrics.Accuracy()
m.update_state(predictions, test_target)
accuracy = m.result().numpy()

print("Prediction accuracy is: ", accuracy)

# Display confusion matrix
print("Confusion Matrix below: ")

tf.math.confusion_matrix(test_target, predictions)
```

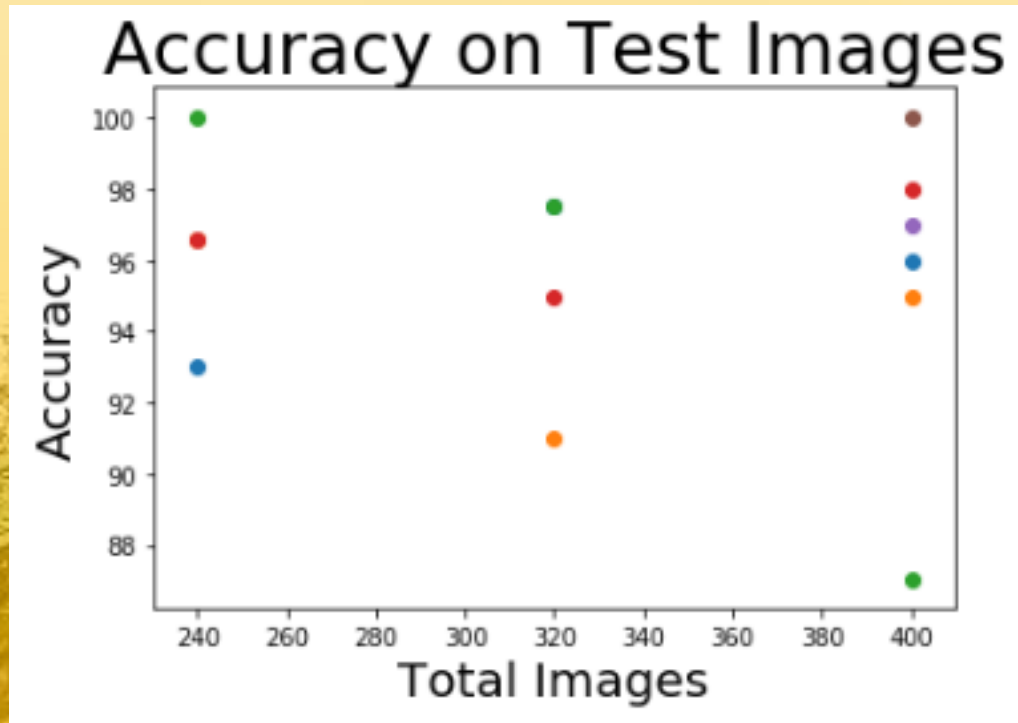
```
Prediction accuracy is: 0.99
Confusion Matrix below:
```

```
Out[17]: <tf.Tensor: shape=(2, 2), dtype=int32, numpy=
array([[51,  1],
       [ 0, 48]], dtype=int32)>
```

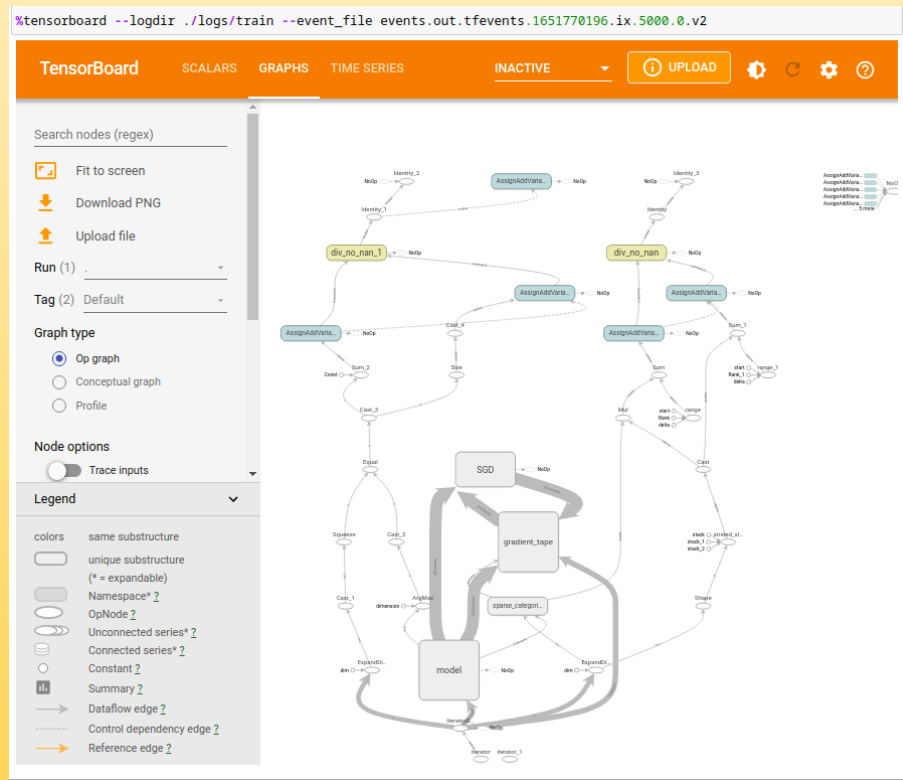
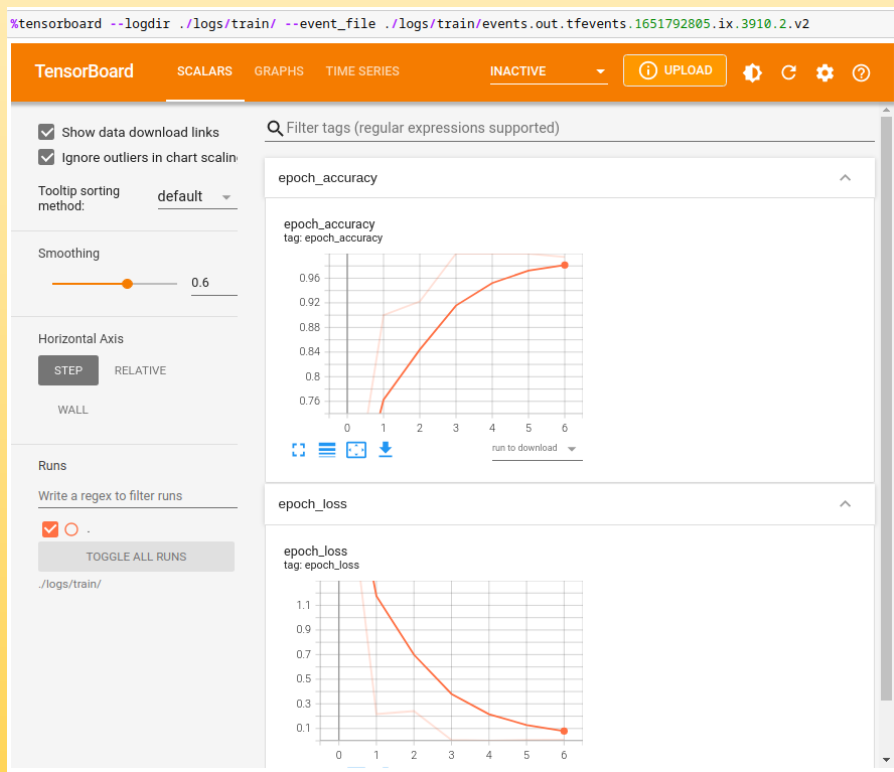
Precision = 1
Recall = 0.98

Column = prediction
Row = target
Starts from 0

***Model achieves 95.7% test accuracy
over 12 runs***



Accuracy improves over time in learning



Model Fine-tuning

1. Unfreeze top 22 layers, trainable weights
 2. Add a dropout layer on the top layer
(don't need dropout if Batch Normalization)
 3. Change optimizer (SGD → Adam)
- **Didn't improve accuracy**

Conclusion

Computer vision with transfer learning is helpful in speed, accuracy and simplicity. And it requires less training data.

The more data we have, the wider distribution the data can be. So the two categories don't have to have a clear cut (i.e. fire and no-fire images look drastically different) to have outstanding prediction accuracy.

Machine Learning can be implemented to detect wildfire to flag the Information for human review and fire management decisions. This saves time and money, and ensures early detection which is important to control the spread of wildfire.

Future Work

Add more data.

No data augmentation.

Implement with higher computational power (e.g. Colab)

References

Aurélien Géron, *Hands-on Machine Learning with Scikit-Learn, Keras & TensorFlow*, 2 Ed. O'Reilly

Abrahams, Hafner, Erwit, Scarpinelli, *TensorFlow for Machine Intelligence*, Blending Edge Press

François Chollet, *Xception: Deep Learning with Depthwise Separable Convolutions*, CVPR, Computer Vision Foundation Open Access version

<https://theaisummer.com/receptive-field/>

<https://www.analyticsvidhya.com/blog/2022/01/convolutional-neural-network-an-overview/>

<https://towardsdatascience.com/4-pre-trained-cnn-models-to-use-for-computer-vision-with-transfer-learning-885cb1b2dfc>

<https://towardsdatascience.com/batch-normalization-the-greatest-breakthrough-in-deep-learning-77e64909d81d>

The background of the slide is a photograph of the University of Toronto's main archway, featuring a central clock face and ornate stone pillars topped with spherical ornaments. The image is overlaid with a semi-transparent yellow filter. The text is centered on the slide.

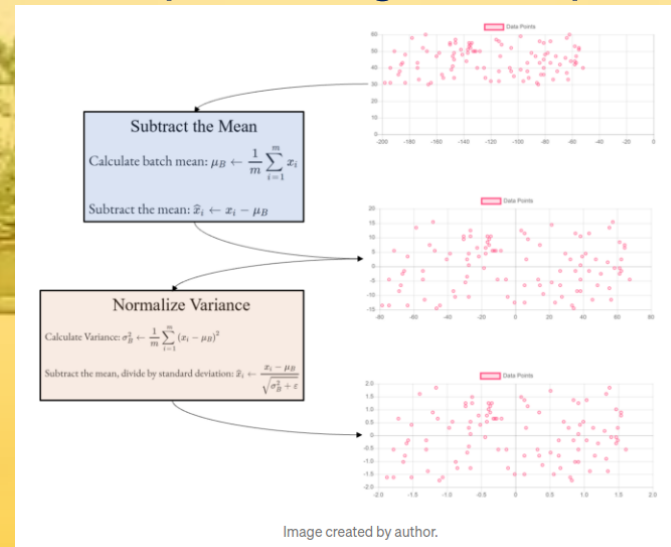
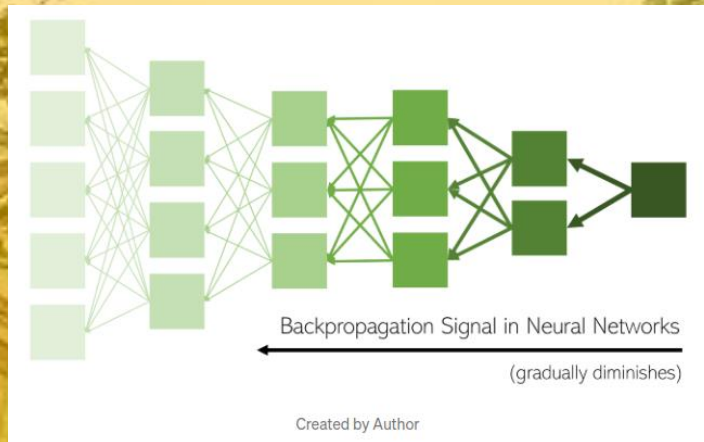
Q&A

Thank you.

Appendix A

Batch Normalization for Deep Learning

Batch Normalization solves gradient vanishing and exploding problems. One of the greatest breakthrough in deep learning development.




Reference

Appendix B:

Chain Rule for Back Propagation

Let's assume a really simply network, with one input, one output, and two hidden layers with a single neuron. Both hidden and output neurons will be sigmoids and the loss will be calculated using cross entropy. Such a network should look like:



Let's define $L1$ as the output of first hidden layer, $L2$ the output of the second, and $L3$ the final output of the network:

$$L1 = \text{sigmoid}(w_1 \cdot x)$$

$$L2 = \text{sigmoid}(w_2 \cdot L1)$$

$$L3 = \text{sigmoid}(w_3 \cdot L2)$$

Finally, the loss of the network will be:

$$\text{loss} = \text{cross_entropy}(L3, y_{\text{expected}})$$

To run one step of gradient decent, we need to calculate the partial derivatives of the loss function with respect of the three weights in the network. We will start from the output layer weights, applying the chain rule:

$$\frac{\partial \text{loss}}{\partial w_3} = \text{cross_entropy}'(L3, y_{\text{expected}}) \cdot \text{sigmoid}'(w_3 \cdot L2) \cdot L2$$

$L2$ is just a constant for this case as it doesn't depend on w_3

To simplify the expression we could define:

$$\text{loss}' = \text{cross_entropy}'(L3, y_{\text{expected}})$$

$$L3' = \text{sigmoid}'(w_3 \cdot L2)$$

The resulting expression for the partial derivative would be:

$$\frac{\partial \text{loss}}{\partial w_3} = \text{loss}' \cdot L3' \cdot L2$$

Now let's calculate the derivative for the second hidden layer weight, w_2 :

$$L2' = \text{sigmoid}'(w_2 \cdot L1)$$

$$\frac{\partial \text{loss}}{\partial w_2} = \text{loss}' \cdot L3' \cdot L2' \cdot L1$$

And finally the derivative for w_1 :

$$L1' = \text{sigmoid}'(w_1 \cdot x)$$

$$\frac{\partial \text{loss}}{\partial w_1} = \text{loss}' \cdot L3' \cdot L2' \cdot L1' \cdot x$$

You should notice a pattern. The derivative on each layer is the product of the derivatives of the layers after it by the output of the layer before. That's the magic of the chain rule and what the algorithm takes advantage of.

We go forward from the inputs calculating the outputs of each hidden layer up to the output layer. Then we start calculating derivatives going backwards through the hidden